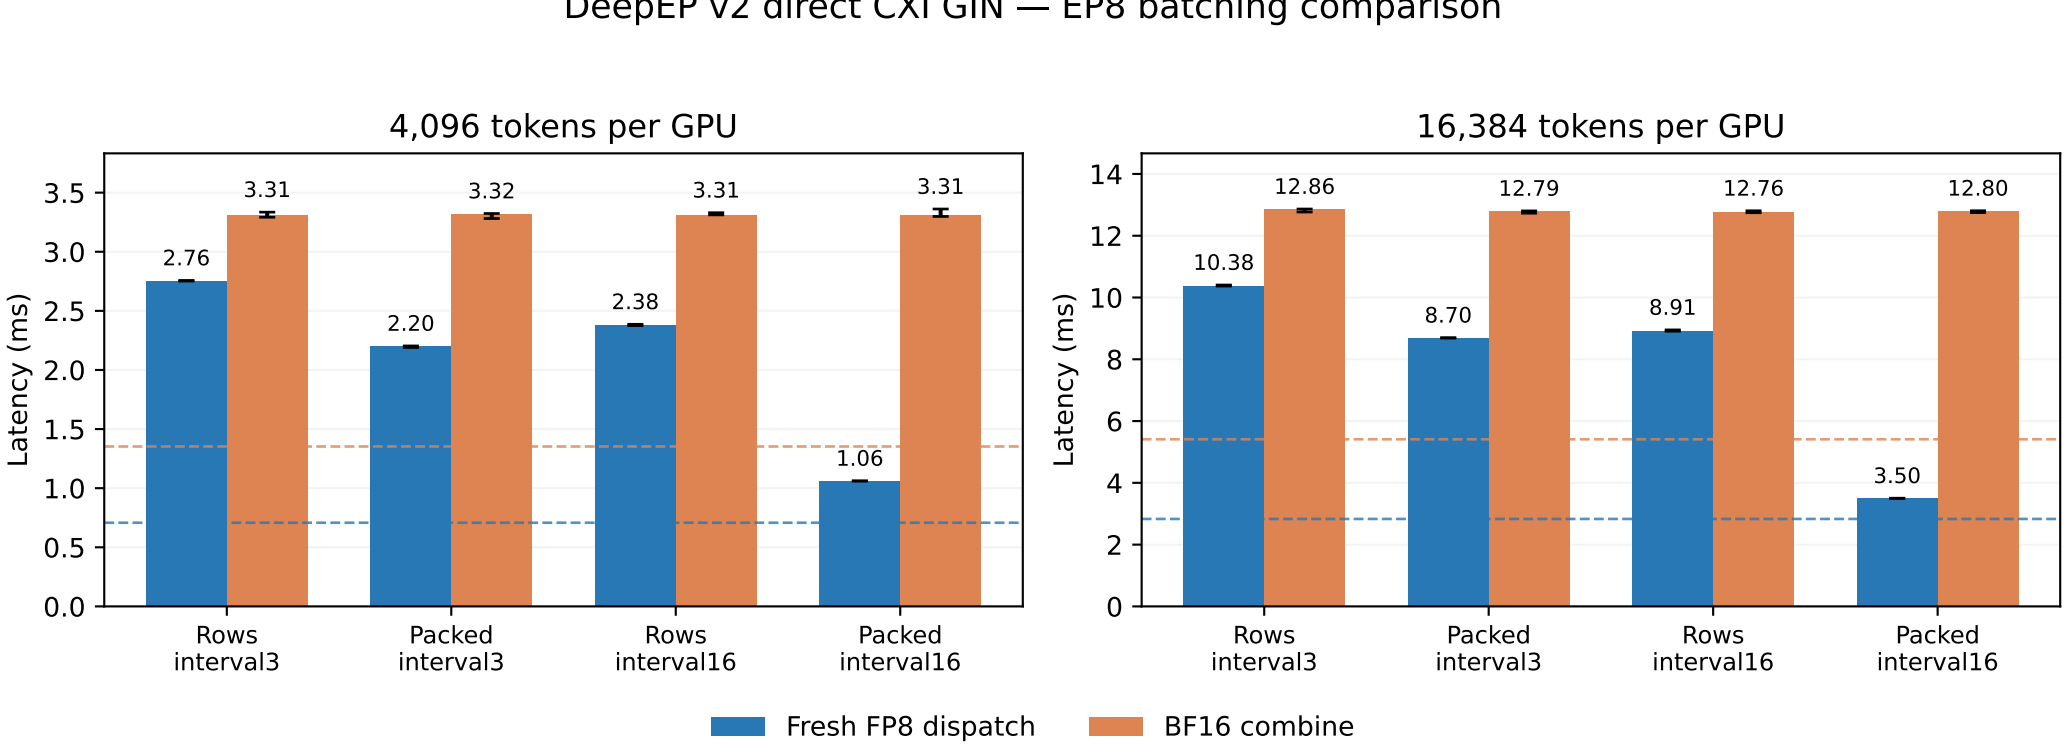
DeepEP v2 direct CXI GIN — EP8 batching comparison
4,096 tokens per GPU
16,384 tokens per GPU
Latency (ms)
2.76
3.31
2.20
3.32
2.38
3.31
1.06
3.31
10.38
12.86
8.70
12.79
8.91
12.76
3.50
12.80
Rows interval3
Packed interval3
Rows interval16
Packed interval16
Fresh FP8 dispatch
BF16 combine
Three repeats; bars: median of rank-max CUDA medians; whiskers: repeat range. Tracing off.
Dashed lines: RDMA payload / 25 GB/s only; relay, packet/control overhead and synchronization excluded.